

Jingfeng Wu

Website: uuujf.github.io
Email: uuujf@berkeley.edu

ACADEMIC EXPERIENCE

University of California, Berkeley

Postdoctoral Fellow

Advisors: Peter Bartlett and Bin Yu

Berkeley, US

2023 - Present

EDUCATION

Johns Hopkins University

Ph.D. in *Computer Science*

Advisor: Vladimir Braverman

Baltimore, US

2019 - 2023

Peking University

M.S. in *Applied Mathematics*

B.S. in *Mathematics & Applied Mathematics*

Beijing, CN

2016 - 2019

2012 - 2016

INDUSTRIAL EXPERIENCE

Google Research

Research Intern

Mentors: Wennan Zhu and Peter Kairuoz

Seattle, US

Summer 2022

Baidu Research

Research Intern

Mentor: Haoyi Xiong

Beijing, CN

Winter 2018

RESEARCH INTERESTS

Deep Learning Theory, Algorithms, Machine Learning, Optimization, Statistical Learning Theory

CONFERENCE PAPERS (asterisk* indicates equal contributions or alphabetical order)

- [1] J. Wu, P. L. Bartlett*, M. Telgarsky*, and B. Yu*. “Benefits of Early Stopping in Gradient Descent for Overparameterized Logistic Regression”. In: *International Conference on Machine Learning (ICML)*. 2025.
- [2] R. Zhang, J. Wu, L. Lin, and P. L. Bartlett. “Minimax Optimal Convergence of Gradient Descent in Logistic Regression via Large and Adaptive Stepsizes”. In: *ICML*. 2025.
- [3] Y. Cai*, K. Zhou*, J. Wu, S. Mei, M. Lindsey, and P. L. Bartlett. “Implicit Bias of Gradient Descent for Non-Homogeneous Deep Networks”. In: *ICML*. 2025.
- [4] H. Zhang, D. Morwani, N. Vyas, J. Wu, D. Zou, U. Ghai, D. P. Foster, and S. M. Kakade. “How Does Critical Batch Size Scale in Pre-training?”. In: *International Conference on Learning Representations (ICLR)*. 2025.
- [5] Y. Cai, J. Wu, S. Mei, M. Lindsey, and P. L. Bartlett. “Large Stepsize Gradient Descent for Non-Homogeneous Two-Layer Networks: Margin Improvement and Fast Optimization”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2024.
- [6] L. Lin, J. Wu, S. M. Kakade, P. L. Bartlett, and J. D. Lee. “Scaling Laws in Linear Regression: Compute, Parameters, and Data”. In: *NeurIPS*. 2024.

- [7] R. Zhang, **J. Wu**, and P. L. Bartlett. “In-Context Learning of a Linear Transformer Block: Benefits of the MLP Component and One-Step GD Initialization”. In: *NeurIPS*. 2024.
- [8] **J. Wu**, P. L. Bartlett*, M. Telgarsky*, and B. Yu*. “Large Stepsize Gradient Descent for Logistic Loss: Non-Monotonicity of the Loss Improves Optimization Efficiency”. In: *Conference on Learning Theory (COLT)*. 2024.
- [9] **J. Wu**, D. Zou, Z. Chen, V. Braverman, Q. Gu, and P. L. Bartlett. “How Many Pretraining Tasks Are Needed for In-Context Learning of Linear Regression?” In: *ICLR*. 2024.
- [10] X. Li, Y. Deng, **J. Wu**, D. Zhou, and Q. Gu. “Risk Bounds of Accelerated SGD for Overparameterized Linear Regression”. In: *ICLR*. 2024.
- [11] **J. Wu**, V. Braverman, and J. D. Lee. “Implicit Bias of Gradient Descent for Logistic Regression at the Edge of Stability”. In: *NeurIPS*. 2023.
- [12] **J. Wu**, W. Zhu, P. Kairouz, and V. Braverman. “Private Federated Frequency Estimation: Adapting to the Hardness of the Instance”. In: *NeurIPS*. 2023.
- [13] H. Li*, **J. Wu***, and V. Braverman. “Fixed Design Analysis of Regularization-Based Continual Learning”. In: *Conference on Lifelong Learning Agents (CoLLAs)*. 2023.
- [14] **J. Wu***, D. Zou*, Z. Chen*, V. Braverman, Q. Gu, and S. M. Kakade. “Finite-Sample Analysis of Learning High-Dimensional Single ReLU Neuron”. In: *ICML*. 2023.
- [15] **J. Wu***, D. Zou*, V. Braverman, Q. Gu, and S. M. Kakade. “The Power and Limitation of Pretraining-Finetuning for Linear Regression under Covariate Shift”. In: *NeurIPS*. 2022.
- [16] D. Zou*, **J. Wu***, V. Braverman, Q. Gu, and S. M. Kakade. “Risk Bounds of Multi-Pass SGD for Least Squares in the Interpolation Regime”. In: *NeurIPS*. 2022.
- [17] **J. Wu***, D. Zou*, V. Braverman, Q. Gu, and S. M. Kakade. “Last Iterate Risk Bounds of SGD with Decaying Stepsize for Overparameterized Linear Regression”. In: *ICML*. 2022.
- [18] **J. Wu**, V. Braverman, and L. F. Yang. “Gap-dependent Unsupervised Exploration for Reinforcement Learning”. In: *International Conference on Artificial Intelligence and Statistics (AISTATS)*. 2022.
- [19] D. Zou*, **J. Wu***, V. Braverman, Q. Gu, D. P. Foster, and S. M. Kakade. “The Benefits of Implicit Regularization from SGD in Least Squares Problems”. In: *NeurIPS*. 2021.
- [20] **J. Wu**, V. Braverman, and L. F. Yang. “Accommodating Picky Customers: Regret Bound and Exploration Complexity for Multi-Objective Reinforcement Learning”. In: *NeurIPS*. 2021.
- [21] H. Li, A. Krishnan, **J. Wu**, S. Kolouri, P. K. Pilly, and V. Braverman. “Lifelong Learning with Sketched Structural Regularization”. In: *Asian Conference on Machine Learning (ACML)*. 2021.
- [22] Z. Yu, C. Hu, **J. Wu**, X. Sun, V. Braverman, M. Chowdhury, Z. Liu, and X. Jin. “Programmable Packet Scheduling with a Single Queue”. In: *ACM Special Interest Group on Data Communication (SIGCOMM)*. 2021.
- [23] D. Zou*, **J. Wu***, V. Braverman, Q. Gu, and S. M. Kakade. “Benign Overfitting of Constant-Stepsize SGD for Linear Regression”. In: *COLT*. 2021.
- [24] **J. Wu**, D. Zou, V. Braverman, and Q. Gu. “Direction Matters: On the Implicit Bias of Stochastic Gradient Descent with Moderate Learning Rate”. In: *ICLR*. 2021.
- [25] J. You, **J. Wu**, X. Jin, and M. Chowdhury. “Ship Compute or Ship Data? Why not Both?” In: *USENIX Symposium on Networked Systems Design and Implementation (NSDI)*. 2021.
- [26] Z. Yu, **J. Wu**, V. Braverman, I. Stoica, and X. Jin. “Twenty Years After: Hierarchical Core-Stateless Fair Queueing.” In: *NSDI*. 2021.
- [27] **J. Wu**, V. Braverman, and L. F. Yang. “Obtaining Adjustable Regularization for Free via Iterate Averaging”. In: *ICML*. 2020.
- [28] **J. Wu**, W. Hu, H. Xiong, J. Huan, V. Braverman, and Z. Zhu. “On the Noisy Gradient Descent that Generalizes as SGD”. In: *ICML*. 2020.
- [29] B. Yu*, **J. Wu***, J. Ma, and Z. Zhu. “Tangent-Normal Adversarial Regularization for Semi-Supervised Learning”. In: *Conference on Computer Vision and Pattern Recognition (CVPR)*. 2019.
- [30] Z. Zhu*, **J. Wu***, B. Yu, L. Wu, and J. Ma. “The Anisotropic Noise in Stochastic Gradient Descent: Its Behavior of Escaping from Sharp Minima and Regularization Effects”. In: *ICML*. 2019.

JOURNAL PAPERS (asterisk* indicates equal contributions or alphabetical order)

- [31] A. Soltoggio, E. Ben-Iwhiwhu, V. Braverman, ..., **J. Wu**, et al. “A collective AI via lifelong learning and sharing at the edge”. In: *Nature Machine Intelligence* (2024).
- [32] D. Zou*, **J. Wu***, V. Braverman, Q. Gu, and S. M. Kakade. “Benign Overfitting of Constant-Stepsize SGD for Linear Regression”. In: *Journal of Machine Learning Research (JMLR)* (2023).

INVITED TALKS

“Reimagining Gradient Descent: Large Stepsize, Oscillation, and Acceleration”

- **MPI & UCLA**, *Math Machine Learning Seminar*, hosted by Guido Montufar June 2025
- **SIAM DS25**, *Dynamical Systems for Machine Learning*, hosted by Molei Tao May 2025
- **UCLA**, *Level Set Meeting*, hosted by Shu Liu and Stanley Osher January 2025
- **Simons Foundation**, *MoDL Annual Meeting*, hosted by Peter Bartlett and Rene Vidal September 2024
- **UC San Diego**, *MoDL Collaboration Meeting*, hosted by Chaoyue Liu et al. May 2024
- **UCLA**, *Computer Science Seminar*, hosted by Quanquan Gu March 2024
- **UC Berkeley**, *Biostatistics Seminar*, hosted by Lexin Li February 2024
- **UC San Diego**, *Group Seminar*, hosted by Mikhail Belkin February 2024

“A Statistical Viewpoint on Implicit Regularization: GD for Logistic Regression”

- **ICTP**, *6th Youth in High-Dimensions Conference*, hosted by Marco Mondelli et al. July 2025
- **UC Berkeley**, *Deep Learning Theory Workshop*, hosted by Peter Bartlett et al. February 2025

“New Insights about SGD: Stepsize, Risk Convergence, and Implicit Regularization”

- **UC Davis**, *Statistics Seminar*, hosted by Xiao Hui Tai October 2023

“Implicit Bias of Gradient Descent for Logistic Regression at the Edge of Stability”

- **TTI-Chicago**, *MoDL Collaboration Meeting*, hosted by Sam Buchanan et al. May 2023

“The Implicit Regularization of SGD in Least Squares and Beyond”

- **Rice University**, *Algorithms and ML Seminar*, hosted by Anastasios Kyrillidis January 2023
- **Princeton University**, *Group Seminar*, hosted by Jason Lee December 2022
- **Georgia Tech**, *Group Seminar*, hosted by Molei Tao November 2022
- **Google Research**, *Learning Theory Seminar*, hosted by Mehryar Mohri August 2022
- **MPI & UCLA**, *Math Machine Learning Seminar*, hosted by Guido Montufar June 2022

PROFESSIONAL SERVICES

Organizer

- Deep Learning Theory Workshop at Simons Institute at UC Berkeley February 2025

Conference Reviewer

- International Conference on Machine Learning (ICML) 2020 - 2025
- Conference on Neural Information Processing Systems (NeurIPS) 2020 - 2025
- International Conference on Learning Representations (ICLR) 2021 - 2025
- International Conference on Artificial Intelligence and Statistics (AISTATS) 2021 - 2023
- Conference on Uncertainty in Artificial Intelligence (UAI) 2023

Conference Program Committee Member

- AAAI Conference on Artificial Intelligence (AAAI)

2021 - 2023

Journal Reviewer

- Journal of Machine Learning Research (JMLR)
- IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)
- Transactions on Machine Learning Research (TMLR)
- SIAM Journal on Mathematics of Data Science (SIMODS)
- IEEE Transactions on Information Theory
- Applied Probability Journals
- Journal of Artificial Intelligence Research (JAIR)

TEACHING

- **Teaching Assistant** for “Machine Learning: Deep Learning”, Johns Hopkins University Spring 2023

HONORS

- **Rising Stars in Data Science**, University of Chicago and UC San Diego 2023
- **MINDS Data Science Fellowship**, Johns Hopkins University Summer 2021
- **Best Reviewers (Top 10%)**, ICML 2021